

Class 8

AUTHOR

Joe Kesler

1.Preparing the data

Data is online here: https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv

```
fna.data <- "WisconsinCancer.csv"
wisc.df <- read.csv(fna.data, row.names = 1)
head(wisc.df)
```

		diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302		M	17.99	10.38	122.80	1001.0
842517		M	20.57	17.77	132.90	1326.0
84300903		M	19.69	21.25	130.00	1203.0
84348301		M	11.42	20.38	77.58	386.1
84358402		M	20.29	14.34	135.10	1297.0
843786		M	12.45	15.70	82.57	477.1
		smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302		0.11840	0.27760	0.3001	0.14710	
842517		0.08474	0.07864	0.0869	0.07017	
84300903		0.10960	0.15990	0.1974	0.12790	
84348301		0.14250	0.28390	0.2414	0.10520	
84358402		0.10030	0.13280	0.1980	0.10430	
843786		0.12780	0.17000	0.1578	0.08089	
		symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302		0.2419		0.07871	1.0950	0.9053
842517		0.1812		0.05667	0.5435	0.7339
84300903		0.2069		0.05999	0.7456	0.7869
84348301		0.2597		0.09744	0.4956	1.1560
84358402		0.1809		0.05883	0.7572	0.7813
843786		0.2087		0.07613	0.3345	0.8902
		area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302		153.40	0.006399	0.04904	0.05373	0.01587
842517		74.08	0.005225	0.01308	0.01860	0.01340
84300903		94.03	0.006150	0.04006	0.03832	0.02058
84348301		27.23	0.009110	0.07458	0.05661	0.01867
84358402		94.44	0.011490	0.02461	0.05688	0.01885
843786		27.19	0.007510	0.03345	0.03672	0.01137
		symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302		0.03003		0.006193	25.38	17.33
842517		0.01389		0.003532	24.99	23.41
84300903		0.02250		0.004571	23.57	25.53
84348301		0.05963		0.009208	14.91	26.50
84358402		0.01756		0.005115	22.54	16.67
843786		0.02165		0.005082	15.47	23.75

	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

Removing the first column (the expert diagnosis)

```
wisc.data <- wisc.df[,-1]
diagnosis <- as.factor(wisc.df[,1])
```

Q1. How many observations in this dataset?

```
dim(wisc.data)
```

[1] 569 30

There are 569 observations.

Q2. How many of the observations have a malignant diagnosis

```
table(wisc.df$diagnosis)
```

B	M
357	212

There are 212 malignant observations.

Q3. How many variables/features in the data are suffixed with _mean?

```
colnames(wisc.data)
```

```
[1] "radius_mean"           "texture_mean"
[3] "perimeter_mean"        "area_mean"
[5] "smoothness_mean"       "compactness_mean"
[7] "concavity_mean"        "concave.points_mean"
[9] "symmetry_mean"         "fractal_dimension_mean"
[11] "radius_se"             "texture_se"
[13] "perimeter_se"          "area_se"
[15] "smoothness_se"         "compactness_se"
[17] "concavity_se"          "concave.points_se"
[19] "symmetry_se"           "fractal_dimension_se"
[21] "radius_worst"          "texture_worst"
[23] "perimeter_worst"       "area_worst"
[25] "smoothness_worst"      "compactness_worst"
[27] "concavity_worst"       "concave.points_worst"
[29] "symmetry_worst"        "fractal_dimension_worst"
```

The function grep could be useful here. How does it work? It searches for matches to argument pattern within each element of a vector.

```
length(grep(pattern = "_mean", colnames(wisc.data) ))
```

```
[1] 10
```

There are 10 variables in the data suffixed with _mean.

2. Principal Component Analysis

Checking the means and standard deviations to see if we need to scale the data:

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01

texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Now that we know we need to scale it, lets run the PCA

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987

```
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005  
Cumulative Proportion  0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997  
PC29      PC30  
Standard deviation    0.02736 0.01153  
Proportion of Variance 0.00002 0.00000  
Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44% of the original variance is captured by the first principal components.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


You get more than 70% of the original variance by PCA 3.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

You get at least 90% of the original variance in the data by PCA 7

Now lets create a simple plot of our pca results:

```
biplot(wisc.pr)
```




Ew > Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is ugly as sin. It's difficult to understand because everything is piled on top of one another and you can't distinguish any individual points from anything else. It just looks like noise.


Lets make a better plot

```
plot(wisc.pr$x[,1],wisc.pr$x[,2], col = diagnosis)
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[,1],wisc.pr$x[,3], col = diagnosis)
```




I notice that benign and malignant are slightly less separate than they were in PCA1 vs PCA2.

Now lets make a nice version of this plot in ggplot 2!

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.1.3

```
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation[,1]
```

<code>radius_mean</code>	<code>texture_mean</code>	<code>perimeter_mean</code>
-0.21890244	-0.10372458	-0.22753729
<code>area_mean</code>	<code>smoothness_mean</code>	<code>compactness_mean</code>
-0.22099499	-0.14258969	-0.23928535
<code>concavity_mean</code>	<code>concave.points_mean</code>	<code>symmetry_mean</code>
-0.25840048	-0.26085376	-0.13816696
<code>fractal_dimension_mean</code>	<code>radius_se</code>	<code>texture_se</code>
-0.06436335	-0.20597878	-0.01742803
<code>perimeter_se</code>	<code>area_se</code>	<code>smoothness_se</code>
-0.21132592	-0.20286964	-0.01453145
<code>compactness_se</code>	<code>concavity_se</code>	<code>concave.points_se</code>
-0.17039345	-0.15358979	-0.18341740
<code>symmetry_se</code>	<code>fractal_dimension_se</code>	<code>radius_worst</code>
-0.04249842	-0.10256832	-0.22799663
<code>texture_worst</code>	<code>perimeter_worst</code>	<code>area_worst</code>
-0.10446933	-0.23663968	-0.22487053
<code>smoothness_worst</code>	<code>compactness_worst</code>	<code>concavity_worst</code>
-0.12795256	-0.21009588	-0.22876753

concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

The component loading loading vector for concave.points.mean is -0.26

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

The minimum number of principal components required to explain 80% of the variance is 4 PC's. You can see this via

```
y <- summary(wisc.pr)
attributes(y)
```

```
$names
[1] "sdev"        "rotation"    "center"      "scale"       "x"
[6] "importance"
```

```
$class
[1] "summary.prcomp"
```

```
sum(y$importance[3, ]<=0.8)
```


```
[1] 4
```

3. Hierarchical Clustering

Lets make a cluster dendrogram that shows our results from the initial table

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled, method = "euclidean")
wisc.hclust <- hclust(data.dist, method = "complete")
plot(wisc.hclust)
abline(h=18, col = "red" ,lty =2)
```

Cluster Dendrogram



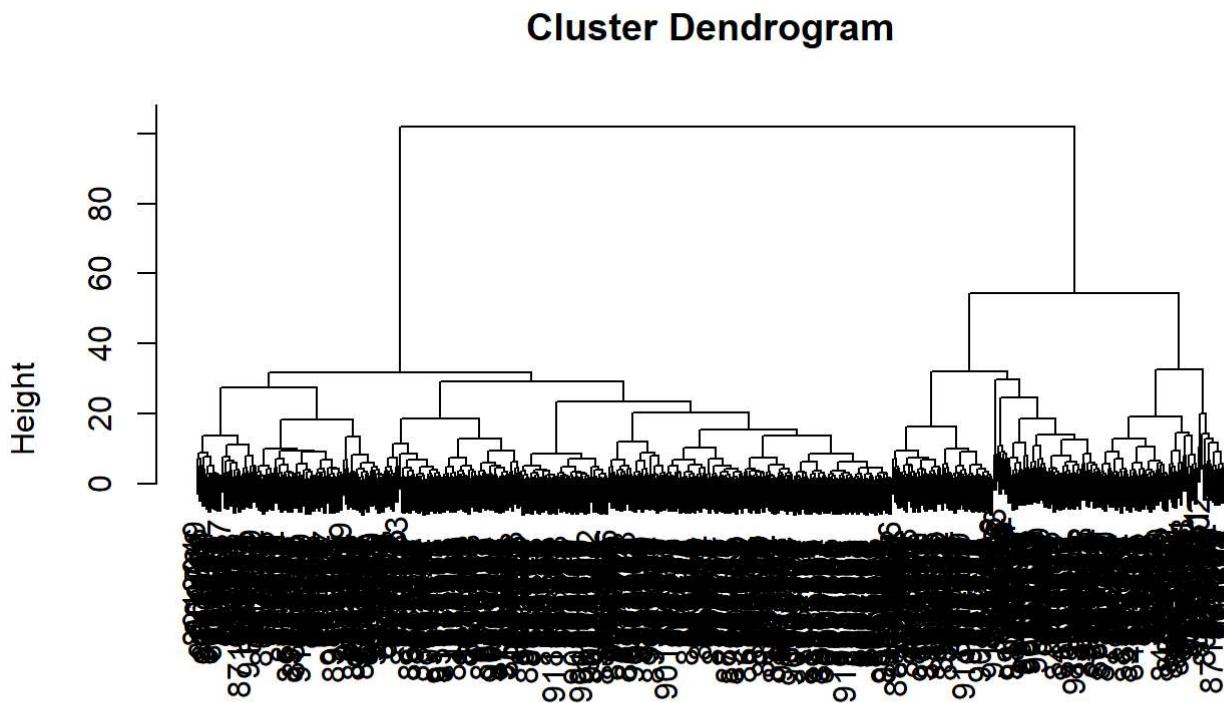
Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

In this case, ward.D2 is my favorite method because it minimizes the variance within clusters.

5. Combining methods

We need to create a model that is a dendrogram of the PCA results.

```
wisc.pr.hclust <- hclust(data.dist, method = "ward.D2")
plot(wisc.pr.hclust)
```



data.dist
hclust (*, "ward.D2")

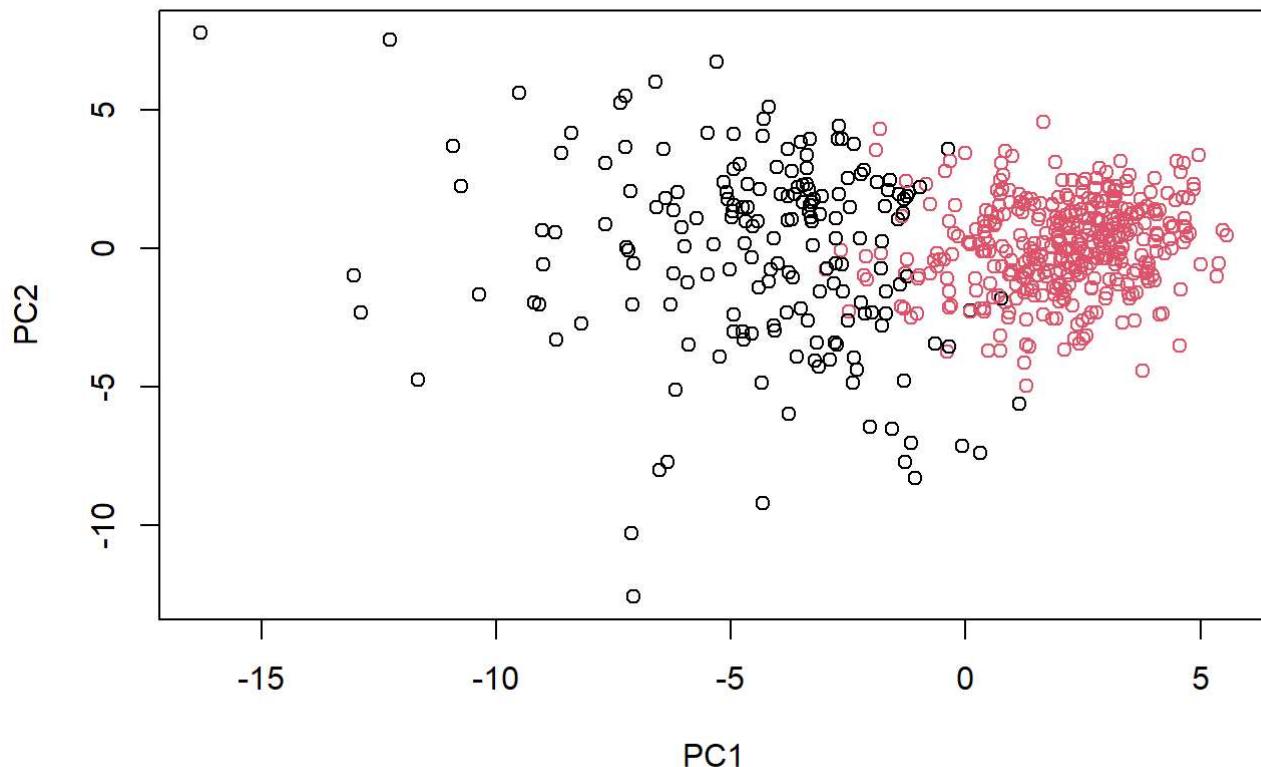
```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
184 385
```


```
table(grps, diagnosis)
```

```
diagnosis
grps   B   M
 1  20 164
 2 337  48
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



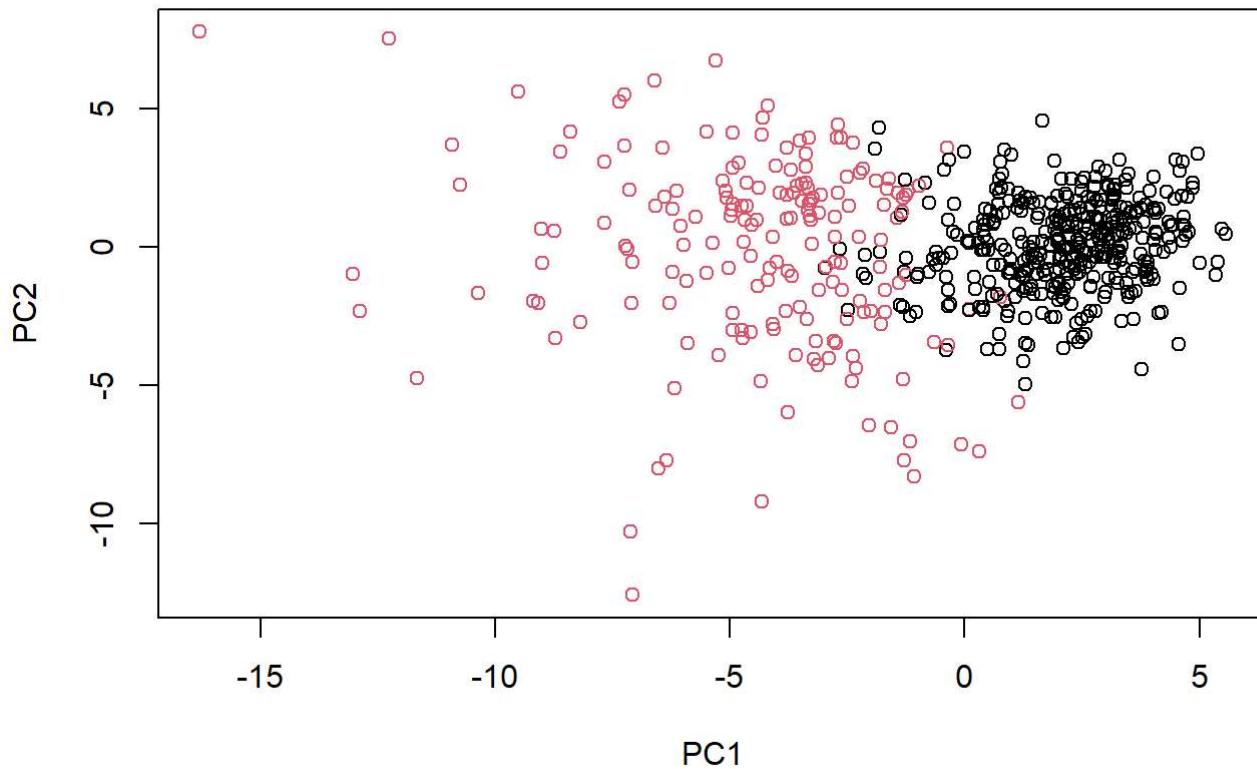
```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method="ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=4)
table(wisc.pr.hclust.clusters, diagnosis)
```

diagnosis		
wisc.pr.hclust.clusters	B	M
1	0	45
2	2	77
3	26	66
4	329	24

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

Four clusters does not split out particularly well because of group "3" which kind of has a lot in both categories. I liked the model with two groups better because it was simpler and easier to generalize as the mostly benign group and mostly malignant group.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
wisc.km <- kmeans(scale(wisc.data), centers= 2, nstart= 20)
table(wisc.km$cluster, diagnosis)
```

```
diagnosis
  B   M
1 14 175
2 343 37
```

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	diagnosis	
	B	M
1	12	86
2	0	59
3	0	3
4	331	39
5	0	20
6	2	0
7	12	0
8	0	2
9	0	2
10	0	1

It seems like they separated the diagnoses well, with the wisc.km cluster seemingly doing better than the PCA model because there is more separation between benign and malignant.

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity?

```
# For the kmeans method:
175/(175+37)
```

```
[1] 0.8254717
```

```
343/(343+37)
```

```
[1] 0.9026316
```

```
# For the PCA method
188/(188+24)
```

```
[1] 0.8867925
```

```
329/(329+24)
```

```
[1] 0.9320113
```

Overall, it seems like PCA method resulted in better of both specificity and sensitivity.